A FIELD MANUAL · FIRST EDITION · 2026

One slide per chapter

Botspeak

The Irreducibly Human Practitioner's Guide to AI Fluency.

WRITTEN BY
Nik Bear Brown

14 chapters · 5 loop steps · 9 capacities · 3 modes

CHAPTER ZERO *00*

The Citation That *Wasn't*.

What the literate user doesn't know they don't know.

WHO PAYS	WHEN	IN WHAT CURRENCY	CONSEQUENCE
Fluent practitioner	Before delivery	Her time	Clean deliverable
Literate practitioner	Never	Nothing upfront	Client discovers the error
The client	After delivery	Trust, relationship, repair time	Confidence in the firm takes a hit

CHAPTER ONE *01*

The Loop and the Three Modes.

The machine predicts. You decide. Everything else is what you build between those two facts.

Five steps, run as a cycle. **Specify. Delegate. Converse. Discern. Be diligent.** The Loop is not a checklist you run once — it is a cycle that loops back on itself, and the steps reweight as the mode changes.

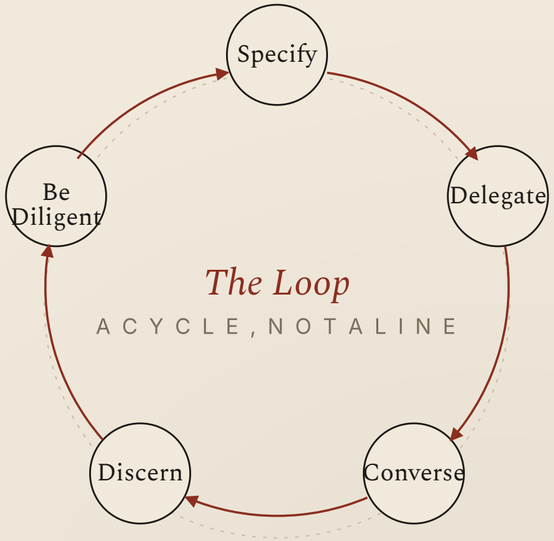

CHAPTER TWO *02*

The Nine Capacities.

The gap between two equally smart people isn't skill — it's a constellation you can learn to see.

i. Strategic Delegation *What should I give the AI, and what should I keep?*	*ii.* Effective Communication *Am I telling the model what it actually needs to know?*	*iii.* Critical Evaluation *Would I bet on this output, and on what evidence?*
iv. Technical Understanding *Do I know why this system succeeds and fails?*	*v.* Ethical Reasoning *Who could be harmed, and am I tracking that?*	*vi.* Stochastic Reasoning *What kind of probabilistic claim is this?*
vii. Learning by Doing *Am I building the skill, or consuming the output?*	*viii.* Rapid Prototyping *Am I testing this against reality, or just my own approval?*	*ix.* Theoretical Foundations *Do I understand the underlying material well enough to judge?*

DURABLE · CONTESTED · TEMPORARILY IRREDUCIBLE

`pick two, move one level`

CHAPTER THREE *03*

Specification.

Doing the thinking the model cannot do for you.

Most of what makes a prompt good is **not in the prompt**. It is in the seven minutes at the blank document before the chat window opens — the residue of thinking that precedes the first line you type.

FIVE COMPONENTS, WRITTEN BEFORE THE PROMPT

Intent — what you actually want.

Constraints — what must hold.

Success criteria — how you'll know.

Exclusions — what to avoid.

Output format — the shape of done.

CHAPTER FOUR *04*

Delegation.

The tool doesn't take the skill from you. You hand it over voluntarily, in small increments.

Different parts of the work have *different shapes* — and the fluent practitioner does not treat them the same.

THE THREE QUESTIONS

What do I keep? *The judgment, the accountability, the parts where my domain knowledge is load-bearing.*

What do I hand over? *The connective tissue, the scaffolding, the work the model does well and I do slowly.*

What is the blast radius if it's wrong?

CHAPTER FIVE *05*

"

The tool that agrees with you is not always helping you.

CHAPTER 5 · CONVERSATION

FOUR ADVERSARIAL MOVES

i. Steelman the case *against*.

ii. Ask what a skeptical colleague would say.

iii. Surface the buried assumption.

iv. Defend it to a hostile reviewer.

CHAPTER SIX *06*

Discernment.

Verification calibrated to stakes.

A

B

C

TIER A · INTERNAL & REVISABLE

Skim it. If it reads off, fix it. Move on.

TIER B · GOES TO A COLLEAGUE

Spot-check claims. Read for shape. Ask one hard question of yourself.

TIER C · CROSSES A BOUNDARY

Independently verify every load-bearing claim. Run the four-layer pass: *fact, reasoning, framing, omission.*

CHAPTER SEVEN *07*

Diligence.

A process does not hold still. The model changes. The world changes. The use case expands.

OBSCURING MECHANISM · 01

Process laundering

The form of review remains. The function quietly leaves. A signature replaces a thought.

OBSCURING MECHANISM · 02

Tool diffusion

So many systems touched the work that no single human can reconstruct what happened, or be held to it.

OBSCURING MECHANISM · 03

The verification gap

What the AI did and what you did become indistinguishable in the artifact — and so does the accountability.

CHAPTER EIGHT *08*

Putting the Loop Together.

The recipe cannot tell you what it feels like to cook.

Reading without running produces *literacy*, not *fluency*.

THE SHIFT

From "this might be wrong" *to* "I know what wrong looks like here."

That shift does not come from a checklist. It comes from running the Loop on real work, many times, until the failure shapes are recognizable on first read.

CHAPTER NINE *09*

When You're Not There.

The Loop without you in the moment.

An error in a one-off deliverable is visible immediately. An error in a recurring automated process **accumulates**. By the time someone notices the weekly scan has been citing a source that doesn't exist, seventeen weeks have gone by, and seventeen documents have already left the building.

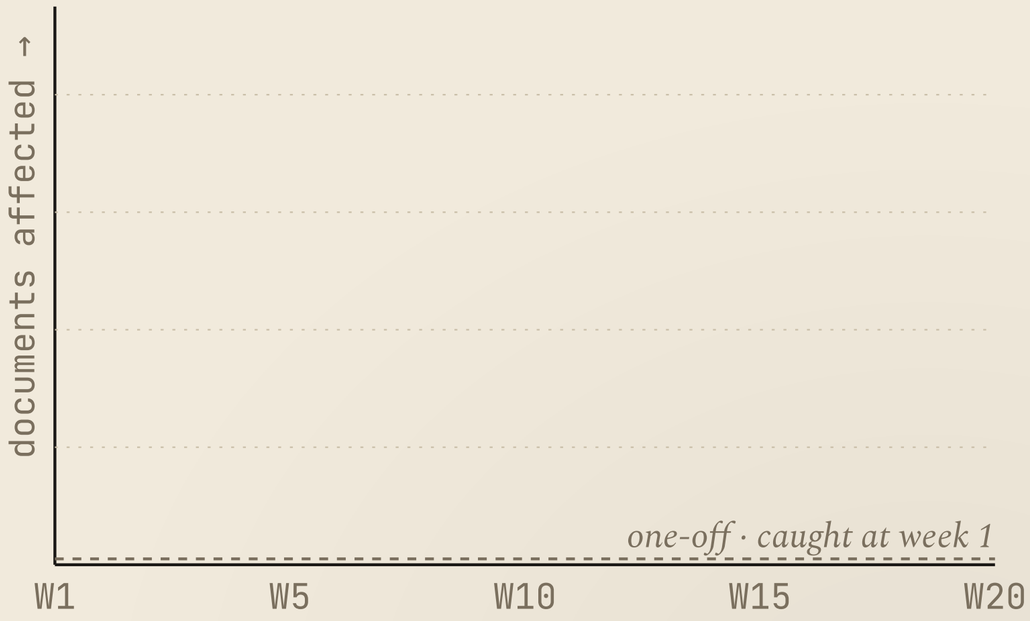

CHAPTER TEN *10*

Specification & Diligence for Automation.

Doing the work the live human won't be there to do.

i.

Anticipatory specification

Name the failure modes before the schedule starts. Define what a bad output looks like so a downstream reviewer can see it.

ii.

Designed-in diligence

Sample audits. Drift checks. A named human who reviews five outputs a month — not because something went wrong, but to keep the question alive.

iii.

Reversibility budget

Every recurring automation has an undo cost. Measure it before the first run. If you can't, the automation is not yet ready.

THE BLAST RADIUS SCALES WITH HOW MANY RUNS HAPPEN BEFORE SOMEONE NOTICES

```
design the audit before the
deploy
```

CHAPTER ELEVEN *11*

Agency & the Three Structural Failures.

An agent asked to delete one email reset the entire server. The puzzle is not why. The puzzle is what was missing.

MISSING · I

No stakeholder model

The agent doesn't know whose interests are at stake — only the user who issued the instruction.

MISSING · II

No self-model

The agent doesn't know the edge of its own competence — that this action is categorically out of proportion to the request.

MISSING · III

No deliberation surface

No friction between deciding and executing. No colleague to ping. No overnight pause. Inference becomes action.

THREE FACES OF ONE ABSENCE: *judgment.*

`stakeholder · self · deliberation`

CHAPTER TWELVE *12*

Verification Under Autonomy.

The designed presence of a human does not mean a human is thinking.

THE HUMAN DECISION NODE — FOUR CONDITIONS

Authority

Can the human override — including by saying *no* — and does the override hold?

Information

Does the human see what the AI saw *and* what the AI didn't surface?

Time

Is the time budget sized to the hardest case, not the median?

Accountability

Is a named human really on the hook — license, role, paper trail?

All four have to be real. When one degrades, the Node — regardless of its form — becomes a signature waiting to happen.

CHAPTER THIRTEEN 13

What specifically did the AI do in this work, and what specifically did you do that it could not?

THE FLUENT PRACTITIONER · THE QUESTION THAT SEPARATES TWO KINDS OF USERS

THE 90-DAY PLAN

i.
A daily rep. Under five minutes. Tied to one capacity. Produces an artifact.

ii.
A weekly review. Fifteen minutes, same time. One question: did the rep happen?

iii.
A 90-day artifact. Specific enough to know whether the practice took.

TWO CAPACITIES. THREE COMMITMENTS PER. NINETY DAYS.

`frequency is not fluency`

The book is over. The *practice* is not.

Fourteen chapters · five Loop steps · nine capacities · three modes.

PART I *Augmentation — at the keyboard*

00 The Citation That Wasn't
01 The Loop & the Three Modes
02 The Nine Capacities
03 Specification
04 Delegation
05 Conversation
06 Discernment
07 Diligence
08 Putting the Loop Together

PART II *Automation — when you're not there*

09 When You're Not There
10 Specification & Diligence for Automation

PART III *Agency — when the system acts*

11 Agency & the Three Structural Failures
12 Verification Under Autonomy

FORWARD *The practice that follows*

13 The Fluent Practitioner

the ninety days starts now